

# Hackathon Report

Manish Nayak

November 21, 2025

## 1 Introduction

The objective is to predict ‘metric score’ (from 0 to 10) for a prompt and its response, for the corresponding metric being used for evaluation. A analysis of the training data shows severe class imbalance, as shown in Table 1.

Table 1: Initial and Modified Score Distribution in Training Data

Initial Distribution		Modified Distribution	
Score	Count	Score	Count
9.0	3123	9.0	3123
10.0	1442	10.0	1443
8.0	259	8.0	259
7.0	95	7.0	95
6.0	45	6.0	46
0.0	13	0.0	13
3.0	7	1.5	11
1.0	6	3.0	10
2.0	5		
4.0	3		
5.0	1		
9.5	1		

## 2 Data Preprocessing and Feature Engineering

- **Text Embedding:** The state-of-the-art `google/embeddinggemma-300m` model from the `SentenceTransformer` library was used to convert all text inputs into 768-dimensional numerical vectors.
- **Feature Creation:** Three distinct embeddings were generated for each data point:
  1. **Metric Embedding:** Retrieved from the pre-computed embeddings file.
  2. **Prompt Embedding:** The `system_prompt` and `user_prompt` were concatenated and encoded to represent the full conversational context.
  3. **Response Embedding:** The ‘response’ text was encoded separately.

- **Data Splitting:** A stratified 80/20 split was performed to create training and validation sets. Stratification on binned score values ensured that the distribution of rare scores was preserved in both sets.

## 3 Modeling Methodology

### 3.1 Model 1: Hierarchical Attention with Optimized Weighted Loss

This model was made to deeply understand the semantic relationships between the prompt, response, and metric.

#### 3.1.1 Architecture

The `HierarchicalAttentionScorer` model utilizes a two-layer cross-attention mechanism:

1. **Prompt-Response Interaction:** The prompt embedding (Query) attends to the response embedding (Key, Value) to produce a "context-aware" prompt that highlights aspects relevant to the given response.
2. **Context-Metric Alignment:** The metric embedding (Query) then attends to this new representation to align the conversation's context with the specific evaluation criterion. The output is fed to a prediction head to generate the final score.

#### 3.1.2 Handling Imbalance

To combat data imbalance, a custom weighted Mean Squared Error (MSE) loss was implemented. This function applies significantly higher penalties for errors on low-score samples. The weights for different score tiers (e.g., "rare", "mid-rare") were treated as hyperparameters.

#### 3.1.3 Optuna Hyperparameter Search

The **Optuna** framework was used to conduct an extensive search (150 trials) to co-optimize both the model's hyperparameters (e.g., number of attention heads, learning rate) and the loss weights. The best parameters found were:

- `num_heads`: 4
- `dropout`: 0.2986
- `lr`: 6.26e-05
- `weight_rare` (scores less than 6): **1001.36**
- `weight_mid_rare` (scores 6-8): **87.92**
- `weight_8`: **11.96**
- `weight_most_common`: 1

The model achieved a best weighted validation RMSE of **9.627**.

## 3.2 Model 2: VQ-VAE, SMOTE-NC, and XGBoost

This model uses a feature-engineering-centric approach to address the imbalance at the data level.

### 3.2.1 Methodology

This was a three-stage pipeline:

1. **Feature Discretization with VQ-VAE:** Three separate Vector Quantized Variational Autoencoders (VQ-VAEs) were trained to learn a discrete "codebook" (size 256) for each embedding type. This converted the high-dimensional, continuous embeddings into single categorical codes, performing a powerful non-linear dimensionality reduction.
2. **Data Balancing with SMOTE-NC:** With the features now categorical, the Synthetic Minority Over-sampling Technique for Nominal and Continuous features (SMOTE-NC) was applied. A tiered strategy synthetically generated new data points for minority classes, with the rarest scores (e.g., less than 6.0) being over-sampled by a factor of 40 times.
3. **XGBoost for Regression:** An XGBoost regressor was trained on the new, balanced dataset of discrete codes. The model's inherent ability to handle categorical features was leveraged.

## 4 Ensembling Strategy and Final Results

The final prediction was generated by ensembling the outputs of both models.

- **Rationale:** Model 1 is a specialist, highly focused on identifying low-scoring samples. Model 2 provides a more generalized perspective from training on a synthetically balanced feature space. By combining their predictions, we leverage the specialized knowledge of Model 1 while regularizing it with the balanced view of Model 2.
- **Method:** A simple average of the predictions from both models was used to produce the final submitted score.

## 5 Conclusion

The Hierarchical Attention Network with an Optuna-tuned weighted loss (Model 1) and the VQ-VAE + SMOTE-NC + XGBoost pipeline (Model 2) represent two distinct methods for handling this challenge.